

# VL Deep Learning for Natural Language Processing

---

## 6. Word Classifiers

*Prof. Dr. Ralf Krestel*

*AG Information Profiling and Retrieval*

# Learning Goals for this Chapter

---

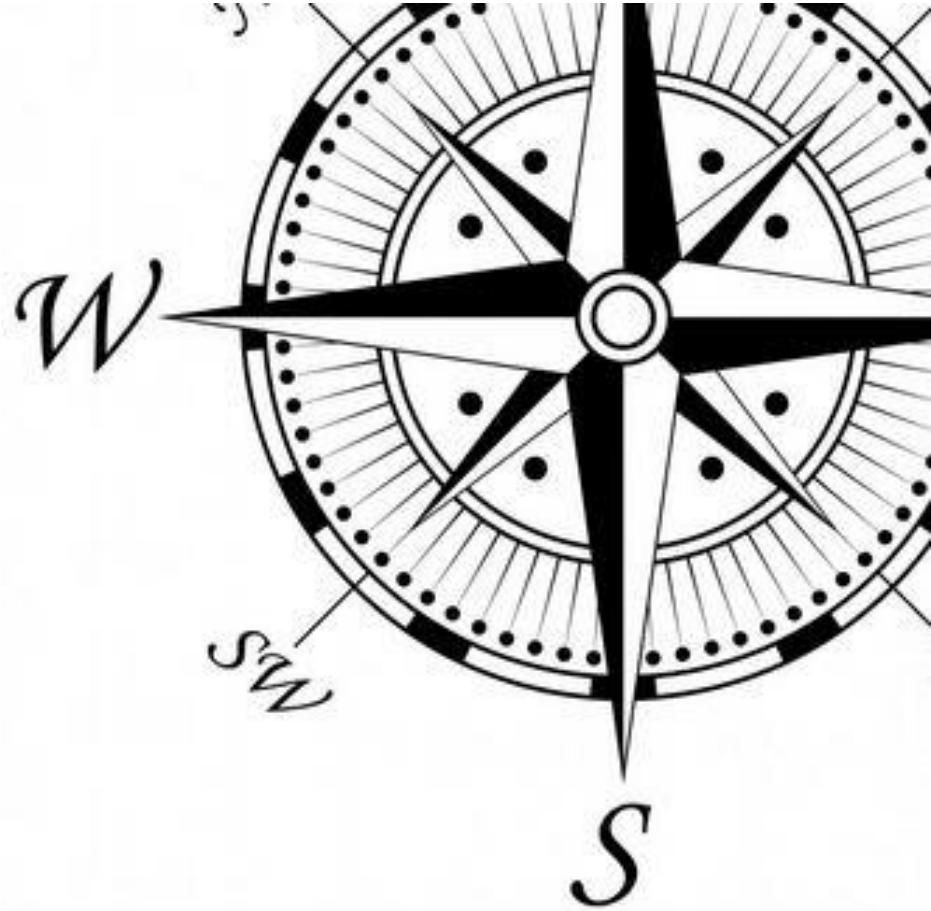


- Know different input representations for text
    - and how/when to use them
  - Develop a simple window classifier
  - Understand the max-margin loss function
    - and know how to employ it
- 
- Relevant chapters
    - P6.1
    - S3 (2019) <https://www.youtube.com/watch?v=PLryWeHPcBs>

# Topics Today

---

1. **Text as Input to Neural Networks**
2. DNN to Classify Words
3. The Max-Margin Loss Function



# What Makes Text Representation Hard?

---



- Noise
  - Spelling mistakes, colloquialisms, slang words, emoticons, ...
- Ambiguity
  - multiple meanings, sarcasm, ...
- Semantic structure
  - order and placement of words can drastically change the meaning, e.g., “not”
- Domain knowledge
  - understanding of the domain and domain jargon necessary
- Multilingualism
  - Mixed-code text, low resource languages

# How to Process Text?

- Token-level
  - Character by character
- Word-level
  - One word at a time
- Sentence-level
  - Or fixed-size context window
- Document-level
  - Fixed-length chunk
  - E.g., first 100 words

Typically: sequential processing



- Recurrent neural networks (LSTM, GRU)
- Seq2seq models
- Transformers

Typically fixed-length, non-sequential



- Feed forward neural networks
- Convolutional neural networks (CNN)
- Generative adversarial networks (GAN)

# How to Represent Text?

Alternatively: multiset of  
1-hot-encoded word vectors



- **Bag-of-words (BOW) model**

- Keep the frequency: word count
- Other weights, e.g., TF-IDF
- Binary: set of words, incident vectors  
(0/1 whether a word occurs at least once)
- Optional: smoothing
- Optional: normalization, stemming, etc.
- Optional: n-grams

- Example

DOC1: *John likes dogs. Mary likes cats. Both play soccer.*

dogs		DOC1
0	<i>John</i>	1
0	<i>likes</i>	2
1	<i>dogs</i>	1
0	<i>Mary</i>	1
0	<i>cats</i>	1
0	<i>Both</i>	1
0	<i>play</i>	1
0	<i>soccer</i>	1
0	...	0
⋮	⋮	⋮
0	<i>Zyzzzyva</i>	0

# How to Represent Text?



- **Topic model**

- Learn a latent (hidden) space
- Based on singular value decomposition (SVD)
  - Latent semantic analysis (LSA) aka latent semantic indexing (LSI)
- Based on probabilistic models
  - Probabilistic LSA (PLSA)
  - Latent Dirichlet allocation (LDA)

		Manual topic titles	
dogs			DOC1
$\begin{bmatrix} 0.8 \\ 0 \\ 0.1 \\ 0 \\ \vdots \\ 0.1 \end{bmatrix}$	<i>nature</i>		$\begin{bmatrix} 0.4 \\ 0.1 \\ 0.4 \\ 0 \\ \vdots \\ 0.1 \end{bmatrix}$
	<i>politics</i>		
	<i>sport</i>		
	...		
	$\vdots$		
	<i>entertain</i>		

- Example

DOC1: *John likes dogs. Mary likes cats. Both play soccer.*

Dimensions in  
latent space

# How to Represent Text?



- **Embedding model**

- Character-level
- Subword-level
- **Word-level**
- Phrase- and sentence-level
- Paragraph-level
- Document-level
- Combinations of levels also possible
  - e.g. character + word for NMT
  - or using lower-levels as additional input

} Most meaningful

Can also be context-dependent

Dimensions in latent space

dogs  
 $\begin{bmatrix} 0.16 \\ 0.21 \\ 0.18 \\ 0.04 \\ \vdots \\ 0.10 \end{bmatrix}$

Meaningful structure along arbitrary directions

DOC1  
 $\begin{bmatrix} 0.05 \\ 0.11 \\ 0.09 \\ 0.02 \\ \vdots \\ 0.04 \end{bmatrix}$

Average or sum of word vectors

Alternatively: Multiset of word vectors

DOC1: *John likes dogs. Mary likes cats. Both play soccer.*



# Most Tasks and Applications

---



- Ordered multiset of words
  - 1-hot-encoded



- Ordered multiset of embeddings
  - Dense word vectors for each word
  - Can be pretrained/fine-tuned

- Sequence of words
  - 1-hot-encoded

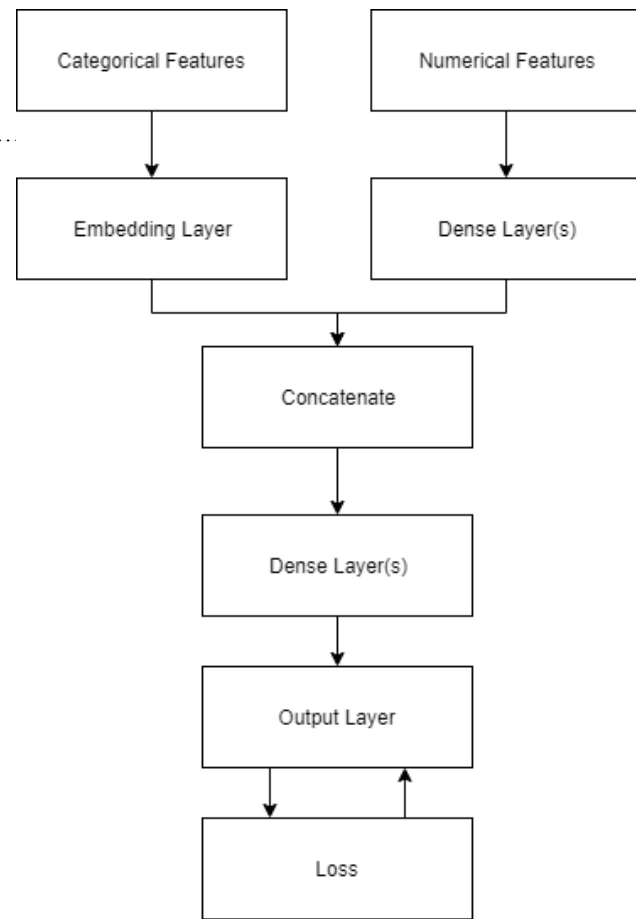


- Sequence of embeddings
  - Dense word vectors for each word
  - Can be pretrained/fine-tuned

# Different Input Data Types

- Categorical
  - Character
  - String (=array of char)
  - Text (indexed strings)
- Numerical
  - Real
  - Integer
- Combining different feature types
  - Concatenating their representations

Objects  
with IDs



# Learning Word Embeddings for Preprocessing

---



- Goal of word2vec: Learn word vectors
  - Iterate through all words in a corpus
  - Predict which words occur in their contexts
  - Thus: Similar context results in similar word vectors!
- But **every** DNN learns a representation of the input data!
  - Word vectors can be learnt for a specific task.
  - Representation is then optimized towards the problem that needs to be solved.
    - For word2vec this problem is to predict the context (meaning).
    - Other nets predict other things, e.g., polarity (sentiment).

# Input Data = Text

---

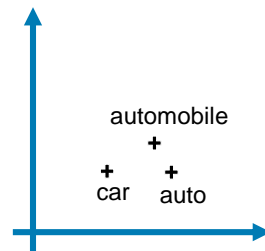


- Deep neural networks need tensors as input
  - One-hot-encoding
    - Bag-of-words
    - Sparsely populated
  - Embedding
    - Dimensionality reduction
    - Densely populated
- Embedding Layer can learn representation **task-specific**
  - First layer in a DNN learns dimensionality reduction / representation
- **Pretrained** Embeddings can be used
  - First layer maps input words to pretrained word vectors

# Pretrained Word Vectors vs. Newly Learned



- Classic machine learning
  - Learning feature weights ( $\mathbb{R}^{Cd}$ )
- Deep Learning
  - Learning feature weights and word vectors ( $\mathbb{R}^{Cd+Vd}$ )
  - $Vd$  is very large
    - Danger of overfitting



Fine-Tuning

- How about a compromise?
  - Load pretrained word vectors and then continue training with current data?
- Problem:
  - Words that occur in the training set move around in the embedding space; words that do not occur in the training set but maybe in the test set stay where they are.

# The Embedding Layer

- Example: DOC2: *John likes dogs. Mary likes cats.*
- Index of words:  $[0,1,2,3,1,4]^T$
- Input matrix I:

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

Why do we need a special embedding layer?  
Why not use a normal fully connected, dense layer?

- Assume a 3-dimensional embedding layer with weight matrix W:

$$\begin{bmatrix} 0.1 & 0.4 & 0.2 \\ 0.4 & 0.3 & 0.2 \\ 0.5 & 0.6 & 0.5 \\ 0.3 & 0.2 & 0.8 \\ 0.6 & 0.6 & 0.1 \end{bmatrix}$$

# The Embedding Layer



- The output would be  $I * W =$

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 0.1 & 0.4 & 0.2 \\ 0.4 & 0.3 & 0.2 \\ 0.5 & 0.6 & 0.5 \\ 0.3 & 0.2 & 0.8 \\ 0.6 & 0.6 & 0.1 \end{bmatrix}$$

$$= \begin{bmatrix} 1 * 0.1 + 0 * 0.4 + 0 * 0.5 + 0 * 0.3 + 0 * 0.6 & 1 * 0.4 + 0 * 0.3 + 0 * 0.6 + 0 * 0.2 + 0 * 0.6 & 1 * 0.2 + 0 * 0.2 + 0 * 0.5 + 0 * 0.8 + 0 * 0.1 \\ \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots \end{bmatrix}$$

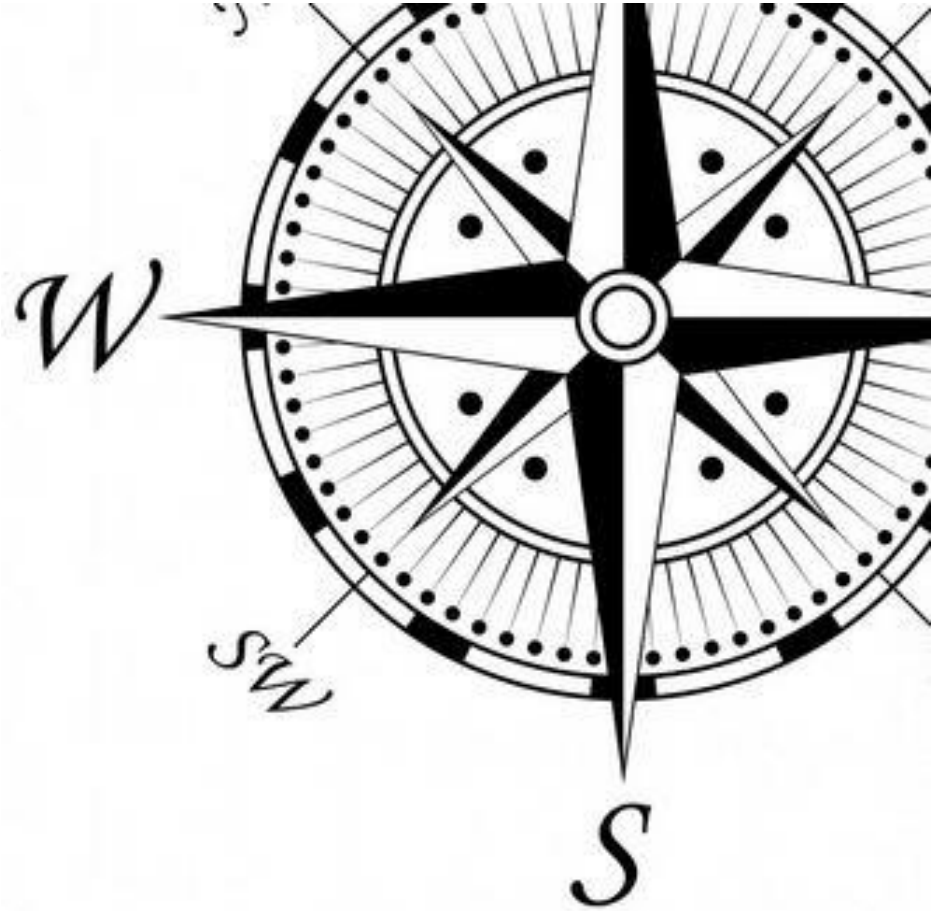
$$= \begin{bmatrix} 0.1 & 0.4 & 0.2 \\ 0.4 & 0.3 & 0.2 \\ 0.5 & 0.6 & 0.5 \\ 0.3 & 0.2 & 0.8 \\ 0.4 & 0.3 & 0.2 \\ 0.6 & 0.6 & 0.1 \end{bmatrix}$$

The embedding layer selects directly the right vectors:  $[0,1,2,3,1,4]^T$

# Topics Today

---

1. Text as Input to Neural Networks
2. **DNN to Classify Words**
3. The Max-Margin Loss Function

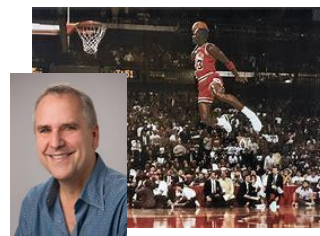
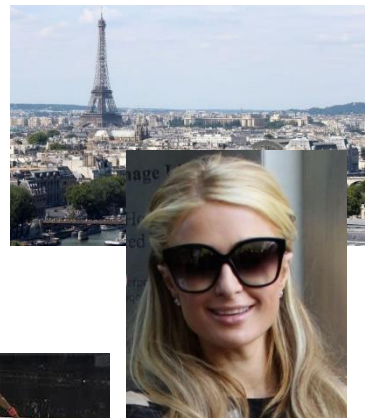




# Word Classification



- Isolated, individual words are rarely classified
  - Meaning of words only clear in context
- Ambiguity can only be resolved in context
  - Autoantonyms: Same words, opposite semantics
    - Overlook, comprise, dust, left, ...
  - Ambiguous named entities
    - Paris
    - Michael Jordan
    - Orange



# Classification with Context

---



- Idea: Classification of a word in the context of its  $k$  neighboring words
- E.g. named entity recognition (NER) can be modeled as word classification with four classes:
  - person, organization, location, none
- Naive approach:
  - Average of all word vectors in the context window.
  - Problem: Position information is lost
- Better:
  - Concatenating all word vectors from the context window

- Softmax classifier
  - Input: the target word to be classified and neighboring words in a context window of size  $k$ , left and right of the target word
  - Output: probability distribution over four named entity classes
- Example: Classification of the word *Paris* as location with context of size two:

...	museums	in	Paris	are	amazing	...
	$\begin{bmatrix} \blacklozenge \\ \blacklozenge \\ \blacklozenge \end{bmatrix}^T$	$\begin{bmatrix} \blacklozenge \\ \blacklozenge \\ \blacklozenge \end{bmatrix}^T$	$\begin{bmatrix} \blacklozenge \\ \blacklozenge \\ \blacklozenge \end{bmatrix}^T$	$\begin{bmatrix} \blacklozenge \\ \blacklozenge \\ \blacklozenge \end{bmatrix}^T$	$\begin{bmatrix} \blacklozenge \\ \blacklozenge \\ \blacklozenge \end{bmatrix}^T$	
$\mathbf{x}_{window} =$	$[\mathbf{x}_{museums}$	$\mathbf{x}_{in}$	$\mathbf{x}_{Paris}$	$\mathbf{x}_{are}$	$\mathbf{x}_{amazing}]^T$	

- The input vector is a column vector in  $\mathbb{R}^{5d}$ 
  - $\mathbf{x}_{window} = [\blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge \blacklozenge]^T$

# Simplest Window Classifier



- Use softmax function!
- With  $x = x_{window}$  we can use softmax on the concatenated word vectors

$$\hat{y}_y = p(y|x) = \frac{\exp(W_y \cdot x)}{\sum_{c=1}^C \exp(W_c \cdot x)}$$

$$W \in \mathbb{R}^{C \times d}$$

- Loss function as usual: cross entropy

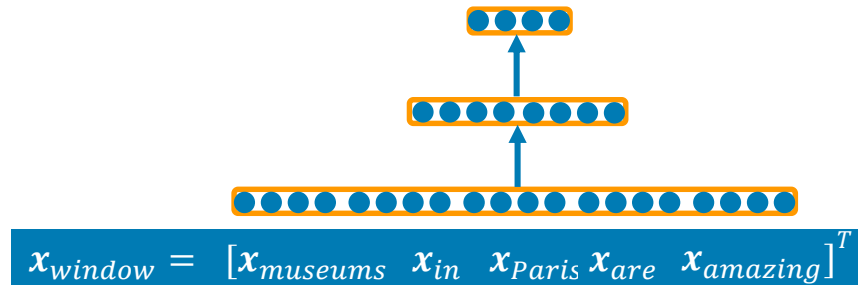
$$J(\theta) = \frac{1}{N} \sum_{i=1}^N -\log \left( \frac{e^{f_{y_i}}}{\sum_{c=1}^C e^{f_c}} \right)$$

$$f_y = W_y \cdot x$$

- How are the word vectors learned?
  - Gradient descend!
  - Analogous to word2vec

# Additional Layer

- Up to now, we only have an input layer (word vectors) and a softmax output layer.
- Softmax can only discriminate linearly in the input space
  - No dependencies of the individual words in the window
- Thus, additional layer with non-linear activation function
  - Example: only if „*museums*“ is the first vector, then the „*in*“ at position two is important



# Single Layer Neural Network

---

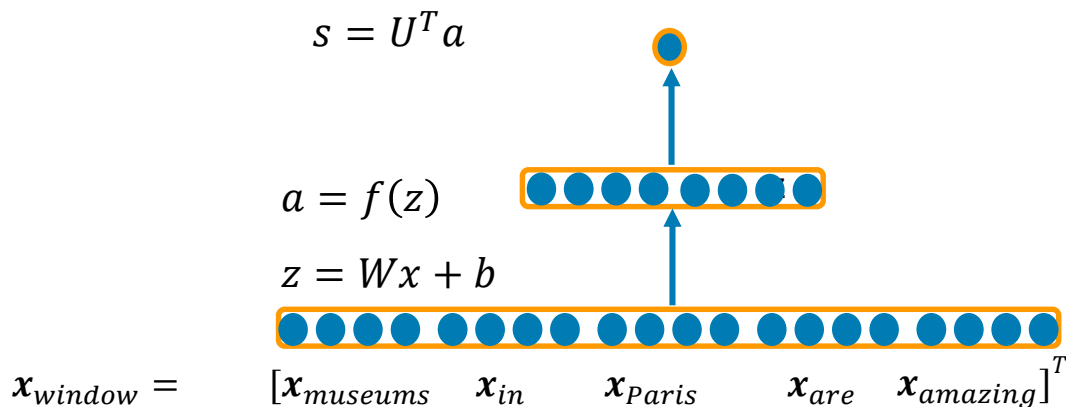


- Input: Concatenated word vectors of a window  $x$
- Hidden layer:  $z = Wx + b$  with activation function  $a = f(z)$
- Output layer:
  - Softmax for probability distribution over multiple classes
    - $p(y|x) = \text{softmax}(Wa)$
  - Easier: non-normalized score for each class
    - $\text{score}(x) = U^T a \in \mathbb{R}$

# Forward Pass



- Score for one window:
  - $s = \text{score}(\text{museums in Paris are amazing})$
  - $s = \mathbf{U}^T f(\mathbf{W}\mathbf{x} + \mathbf{b})$     $\mathbf{x} \in \mathbb{R}^{20 \times 1}$ ,  $\mathbf{W} \in \mathbb{R}^{8 \times 20}$ ,  $\mathbf{U} \in \mathbb{R}^{8 \times 1}$





- How does the network architecture look like if softmax over four classes should be used for the output layer?
- Which effect do the window size and the dimensionality of the word vectors have?
  - On runtime,
  - memory,
  - Quality of results?

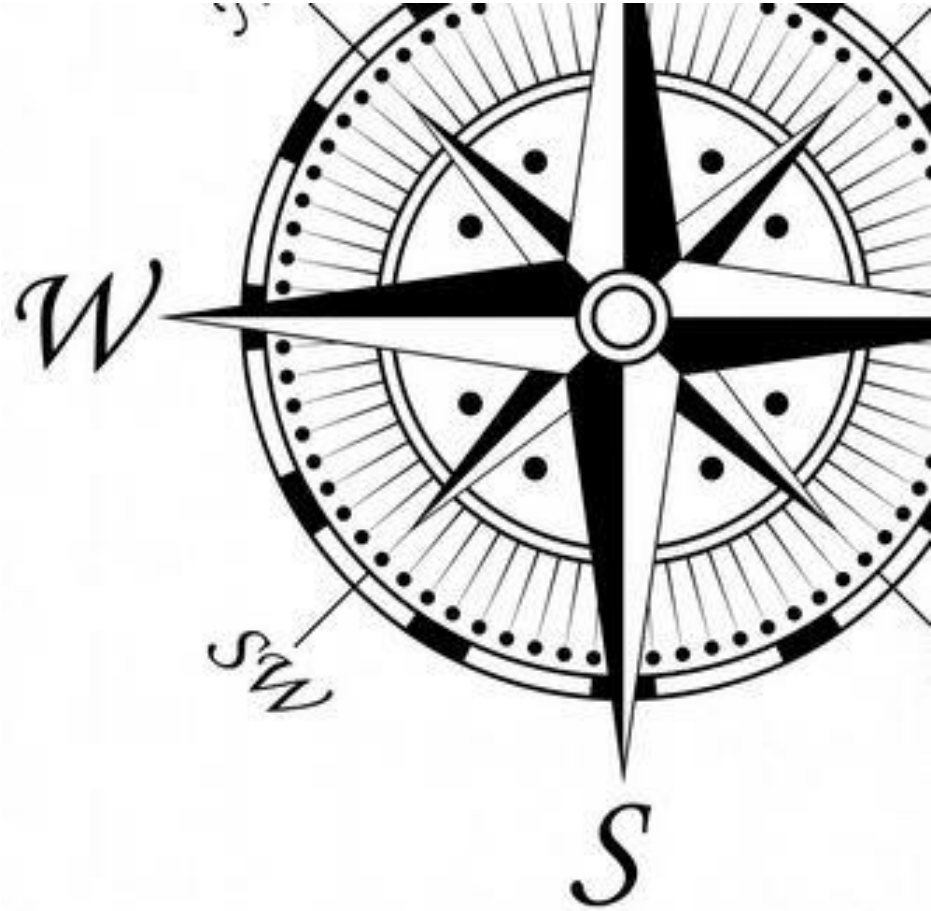





# Topics Today

---

1. Text as Input to Neural Networks
2. DNN to Classify Words
3. **The Max-Margin Loss Function**



# Max-Margin Loss Function Idea

- Many windows to train
  - Most of them don't have a location in the center
    - Those are called „corrupt“ or negative samples
- Example
  - $s = \text{score}(\text{museums in Paris are amazing})$
  - $s_c = \text{score}(\text{Not all museums in Paris})$
- Score of positive examples should be high; score of negative examples should be low
- Target function 
  - $J = \max(0, 1 - s + s_c)$
  - Not differentiable but continuous
    - SGD

# For One Training Sample

---

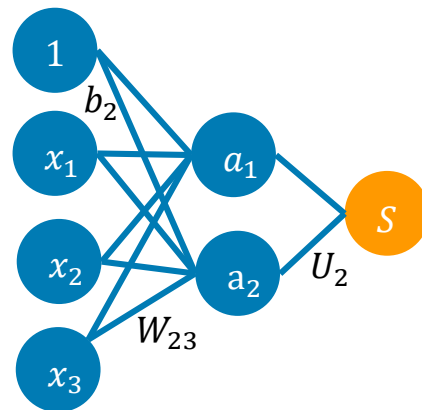


- $J = \max(0, 1 - s + s_c)$
- For each positive window, i.e. with a location in the center
  - Sample  $m$  negative windows
  - Sum over these training windows
- The score for each positive window should be at least 1.0 higher than the highest score of the negative windows

# Training with Backpropagation

- $J = \max(0, 1 - s + s_c)$ 
  - Assumption at start  $J > 0$
- $s = U^T f(Wx + b)$
- $s_c = U^T f(Wx_c + b)$
- Derivatives of  $s$  and  $s_c$  with respect to all involved variables:  $U, W, b, x$
- Gradient with Respect to  $U$

$$\frac{\partial s}{\partial U} = \frac{\partial}{\partial U} U^T a = a$$



# Gradient with Respect to W

- $\frac{\partial s}{\partial W} = \frac{\partial}{\partial W} U^T a = \frac{\partial}{\partial W} U^T f(z) = \frac{\partial}{\partial W} U^T f(Wx + b)$
- Derivative of single weight  $W_{ij}$

$$\frac{\partial s}{\partial W_{ij}} = \frac{\partial}{\partial W_{ij}} U_i a_i$$

$$= U_i \frac{\partial}{\partial W_{ij}} a_i$$

$$= U_i \frac{\partial a_i}{\partial z_i} \frac{\partial z_i}{\partial W_{ij}} = U_i \frac{\partial f(z_i)}{\partial z_i} \frac{\partial z_i}{\partial W_{ij}}$$

$$= U_i f'(z_i) \frac{\partial z_i}{\partial W_{ij}} = U_i f'(z_i) \frac{\partial W_{i \cdot} x + b_i}{\partial W_{ij}}$$

$$= U_i f'(z_i) \frac{\partial}{\partial W_{ij}} \sum_k W_{ik} x_k$$

$W_{ij}$  only appears within  $a_i$

$W_{ij}$  only appears together with  $x_j$

$$= U_i f'(z_i) x_j$$

with  $f'(z) = f(z)(1 - f(z))$  for the logistic function

Local error signal  $\delta_i$

Local input signal  $x_j$

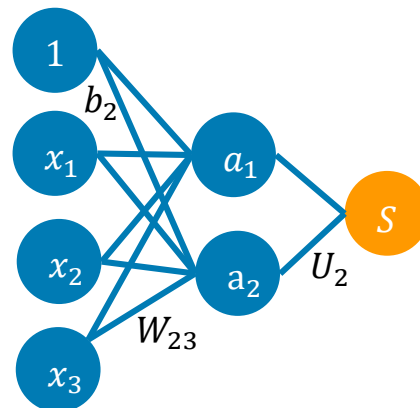
SS 2024

VL DL4NLP

29

$$z_i = W_{i \cdot} x + b_i = \sum_k W_{ik} x_k + b_i$$

$$a_i = f(z_i)$$



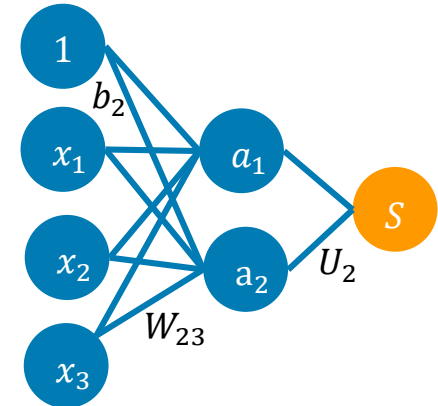
# Gradient with Respect to $W$

- $\frac{\partial s}{\partial W} = \frac{\partial}{\partial W} U^T a = \frac{\partial}{\partial W} U^T f(z) = \frac{\partial}{\partial W} U^T f(Wx + b)$
  - From single weight  $W_{ij}$  to full  $W$ :  

$$\frac{\partial s}{\partial W_{ij}} = U_i f'(z_i) x_j = \delta_i x_j$$
  - We want all combinations  $i=1,2$  and  $j=1,2,3$ 
    - We need all, i.e., complete matrix  $W$
  - Solution:
    - Dyadic product (outer product), often denoted  $\otimes$
    - In contrast to scalar product (inner product)
- $\frac{\partial s}{\partial W} = \delta x^T$  with  $\delta \in R^{2 \times 1}$  is the error signal coming from each activation  $a$

$$z_i = W_i \cdot x + b_i = \sum_k W_{ik} x_k + b_i$$

$$a_i = f(z_i)$$

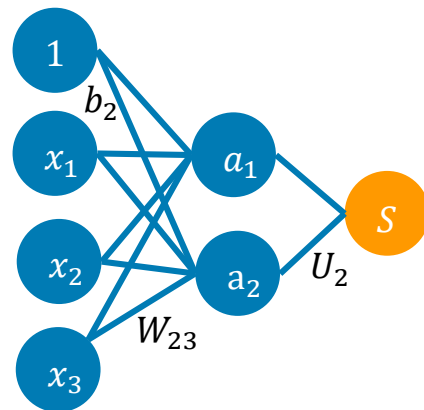


Dimensions of  $\frac{\partial s}{\partial W}$ ?

# Gradient with Respect to b



- $J = \max(0, 1 - s + s_c)$
- $s = U^T f(Wx + b)$
- $s_c = U^T f(Wx_c + b)$
- $\frac{\partial s}{\partial U} = \frac{\partial}{\partial U} U^T a = a$
- $\frac{\partial s}{\partial W} = \delta x^T$
- $\frac{\partial s}{\partial b_i} = U_i f'(z_i) \frac{\partial W_{i \cdot} x + b_i}{\partial b_i} = \delta_i$
- $\frac{\partial s}{\partial b} = \delta$



# Gradient with Respect to $x$

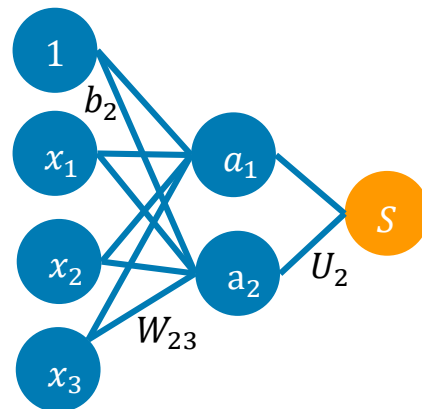
- $J = \max(0, 1 - s + s_c)$
- $s = U^T f(Wx + b)$
- $s_c = U^T f(Wx_c + b)$

- $\frac{\partial s}{\partial x_j} = \sum_{i=1}^2 \frac{\partial s}{\partial a_i} \frac{\partial a_i}{\partial x_j}$
- $= \sum_{i=1}^2 \frac{\partial U^T a}{\partial a_i} \frac{\partial a_i}{\partial x_j}$
- $= \sum_{i=1}^2 U_i \frac{\partial f(W_i \cdot x + b)}{\partial x_j}$
- $= \sum_{i=1}^2 U_i f'(z_i) \frac{\partial W_i \cdot x}{\partial x_j}$
- $= \sum_{i=1}^2 \delta_i W_{ij}$
- $= W_{.j}^T \delta$

$$\frac{\partial s}{\partial x} = W^T \delta$$

Dimensions of  $\frac{\partial s}{\partial x}$ ?

Dimensions of  $\frac{\partial s}{\partial x_j}$ ?





# All Gradients



- $J = \max(0, 1 - s + s_c)$
- $s = U^T f(Wx + b)$
- $s_c = U^T f(Wx_c + b)$
- Derivatives of  $s$  and  $s_c$  with respect to all involved variables:  $U, W, b, x$
- $\frac{\partial s}{\partial U} = \frac{\partial}{\partial U} U^T a = a$
- $\frac{\partial s}{\partial W} = \delta x^T$
- $\frac{\partial s}{\partial b} = \delta$
- $\frac{\partial s}{\partial x} = W^T \delta$
- From  $s$  to  $J$ : indicator function:  $\mathbb{I}[1 - s + s_c > 0] = \begin{cases} 1 & \text{if } 1 - s + s_c > 0 \\ 0 & \text{else} \end{cases}$
- E.g. gradient of  $J$  with respect to  $U$ :  $\frac{\partial J}{\partial U} = \mathbb{I}[1 - s + s_c > 0](-a + a_c)$

# Learning Goals for this Chapter

---



- Know different input representations for text
    - and how/when to use them
  - Develop a simple window classifier
  - Understand the max-margin loss function
    - and know how to employ it
- 
- Relevant chapters
    - P6.1
    - S3 (2019) <https://www.youtube.com/watch?v=PLryWeHPcBs>

- Li, J., Sun, A., Han, J., & Li, C. (2020). A survey on deep learning for named entity recognition. *IEEE Transactions on Knowledge and Data Engineering*, 34(1), 50-70.
  - <https://arxiv.org/pdf/1812.09449.pdf>